




Everyone is building amazing things with Claude right now.

And I don't understand a word of it.






Come on, I'll explain.



Let's start with
CLAUDE.md. Why
does everyone have one?

CLAUDE.md is basically a user manual
that you hand to Claude.




It says how your project works, which
rules apply - and what Claude must
never do.



So I don't have to explain the same thing every session?

Exactly. Claude reads the file automatically as soon as you start inside the project.

There are even three levels: globally in `~/.claude/`, in the project folder and per subfolder. `/init` writes you a first draft.



Use this style

Run these tests




Never touch this folder

Claude sticks to the project rules

Okay. So what are skills, then?

A skill teaches Claude **how** a particular job is done properly.


Instead of explaining your whole procedure every time, you turn it into a reusable skill.



Give me an example.

Say you want the same analysis report every week.


Your instructions, your procedure, your examples – that becomes a skill you simply call.



And those SKILL.md files everyone downloads from GitHub – they're just instructions?

Pretty much. A skill is a folder with a `SKILL.md` in it, plus scripts, templates and reference files.




Claude first reads only the name and description – the rest only when it really needs it. You can also call it yourself with `/skill-name`.



And plugins? Aren't those just skills as well?

Not quite. A skill is **one** ability.


A plugin is more of a **bundle**: skills + hooks + subagents + MCP servers + custom commands in one go. You install it from a marketplace with `/plugin`.



Hold on – MCP servers.
What are those?

MCP is basically how Claude connects
to things **outside** of Claude.


One standard plug for other people's
programs and services.



So I don't have to paste everything in by hand any more ...

Claude connects on its own, fetches the information – and may act there too, if you allow it.


`/mcp` shows you what is currently connected.



Right. And hooks? That sounds even more confusing.

Hooks are different because **Claude** **doesn't** decide whether they happen.


They fire automatically as soon as a certain event occurs. Set up in `settings.json` or through `/hooks`.



For example?

A hook can run your formatter after every file change – or stop Claude **before** a dangerous command runs.

Typical events: `PreToolUse`,
`PostToolUse`, `UserPromptSubmit`,
`SessionStart`, `Stop`.



Claude changes a file



HOOK runs automatically

Format



Test

Check

Report

STOP


Stops before dangerous commands



What about subagents?
Are those literally
several Claudes?

They are independent workers Claude can hand clearly defined jobs to – each with its own context window and its own model.



One researches, one reviews code, one tests. Then they report back to the main Claude. Managed through `/agents`.



And what are all those
/something commands
everyone uses?

The slash is Claude Code's shortcut
menu.

There are built-in ones like `/clear`
and `/compact` - and skills bring their
own commands that you can call
yourself.



`/clear`

start over

`/compact`

boil the context down

`/context`

what eats the space?

`/rewind`


rewind

`/model`

switch model

`/usage`



check your usage



Is Claude allowed to do anything at all on my machine?

No. Out of the box it may read – for anything else it asks first and shows you exactly what it intends to do.

How strict it is you switch with `Shift+Tab`. In plan mode Claude first looks at everything and puts a proposal in front of you before anything happens.



Manual

reads on its own, asks before every change

Accept edits

may touch files, everything else still asks

Plan

only looks and proposes – changes nothing

Automatic



works straight through, with a safety net behind it

switch with Shift + Tab

Are there shortcuts worth knowing?

Four characters save most of the typing – and two keys are the emergency exit.

Remember `Esc` above all: you don't have to wait until Claude is done. Stop it, put it right, carry on.



`@file`

pull a file into the conversation

`!command`

run something yourself

`/`

open the command menu

`Shift+Tab`

switch permission mode

`Esc`

stop Claude right now




`Esc Esc`

rewind

Last question: what exactly is this context window?

That is Claude's working memory for the conversation you are having.




Your input, Claude's answers, files it has read and tool results all take up room in it.



So when people say
"Claude lost the plot" ...

... then usually so much has piled up
that it no longer all fits into the
working memory at once.

Claude Code then summarises
automatically. `/context` shows you
what is taking up the space, `/clear`
gives you a clean start.




What Claude currently has in
mind

And if Claude messes up?
Is everything lost?

No – Claude Code sets checkpoints automatically before it changes files.


`[Esc]` twice or `/rewind`, and you roll back: just the conversation, just the files, or both.



this is where it went wrong




`/rewind – back to state B`



Does Claude remember anything beyond a single session?

Yes – that is what memory is for: small note files Claude writes itself and reads again the next time it starts.

Your preferences, where a project stands, decisions that keep coming back. Unlike CLAUDE.md, Claude maintains this itself.




survives from session to session

And all of this only runs in the terminal?

Not for a long time now. Claude Code comes as a terminal command, as a desktop app, in the browser at `claude.ai/code` and as an extension for VS Code and JetBrains.

On top of that, agents that keep working in the background or in the cloud while you do something else.





Terminal

Desktop app

Browser

VS Code / JetBrains





the same session,
everywhere

Is there a version of all this to click on?

Yes. The Claude desktop app has three tabs: **Chat** is the conversation, **Cowork** does office work, **Code** is Claude Code with a user interface.

It is available for Mac and Windows, and as a beta for Linux.



Chat

Cowork

Code

Talking

questions, drafts, ideas

Delegating

documents and folders




Developing

source code and Git

What can the app do that the terminal can't?

Several sessions side by side – each in its own Git worktree. That way they don't get in each other's way, even when they work on the same spot.

Plus editor, terminal, preview and the review of the changes in the same window. You can even kick off a session from your phone.





everything in one window

So what is Cowork, then?

The same machinery as Claude Code – only in the ordinary Claude app instead of the terminal, with nothing to set up.

You share folders, describe the goal, and Claude gets going: spreadsheets with real formulas, presentations, finished documents.



Your folders

Your accounts



COWORK



Spreadsheet



Presentation

Report

So what does it actually do?

Turn a folder full of receipt photos into an expense report. Write a report from a pile of notes. Tidy up a filing system. If need be it will drive Chrome for that – clicking, typing, filling in forms.

How much it may decide on its own is up to you.



Receipts → expense report

Notes → report

Tidy up the filing

and in the browser if that's what it takes

Asks every time
you approve everything



Decides for itself
with a safety net

Just gets on with it
without asking

When do I use which?

The question is always how far Claude may go into your things. Cowork gets individual folders, runs sealed off and needs no setting up.

Claude Code gets the whole machine. Only here can you teach it tools of your own – and only here does it keep going when you are away.



1 · In the conversation nothing to set up – you upload, you download

2 · Cowork desktop app: Claude works in your folders



3 · Claude Code you are here: the whole machine, your own tools, runs without you

getting things done → Cowork · building tools that run by themselves → Code

And what if something should run regularly on its own?

Then you create a routine: task, project and accounts saved once – after that it runs without you.

And it does so in the cloud. Your machine may be switched off. You create it with `/schedule`, in the app or in the browser.



`/schedule review the new pull requests daily at 9am`



Task

Project

Accounts



runs in the cloud – machine off

And what sets a routine off?



Three things, alone or combined: a schedule, a call from outside, or an event on GitHub.

The shortest interval is one hour.

One-off appointments work too -

/schedule in two weeks ...

The whole thing is still in preview.



Schedule
hourly to weekly

A call
from your own tools

GitHub
new pull request



Routine runs





the result is waiting in the morning

And if it's only meant
for the next hour?

Then use `/loop`. It repeats a task
inside the open session – on your
machine, with your files.

Leave out the interval and Claude picks
it itself: short while something is
happening, longer when things are
quiet. `[Esc]` ends the loop.



`/loop 5m check whether the build has finished`



with an interval
every 5 minutes

without one
Claude decides



with no task at all
it tidies up by itself

runs only while the session is open – and ends after seven days

That sounds almost like a routine.

The difference is where it runs. The loop needs your session open and your machine switched on – in return it sees your files.


The routine runs in the cloud, without you. Rule of thumb: loop for this afternoon, routine for every Monday.



And which model is actually doing the work?

As of August 2026: **Opus 5** for the heavy lifting, **Sonnet 5** for everyday work, **Haiku 4.5** for fast and cheap.

Switch with `/model`. With `/fast` Opus answers more quickly – it is not swapped for a smaller model in the process.



Opus 5

the heavy
jobs

1M context

Sonnet 5

everyday
work


1M context

Haiku 4.5

fast & cheap

200K context



switch with `/model`



And how do I best put it?

State the goal, not the individual steps
– Claude finds the way itself. And say how you can tell when it is done.

And give the context: who it is for, what for, and what must not happen under any circumstances. The clearer the task, the fewer rounds.



rather not

“Make the table look nicer.”





better

“The table goes to the board. Numbers right-aligned, totals row in bold, no colours. Check at the end that the totals add up.”

Can I actually rely on it?

Usually yes – but Claude can also be convincingly wrong. Have it show you what it did instead of just believing it.

And whatever can be checked, let Claude check itself: run the tests, redo the arithmetic, name the sources.



Look at the changes
not just the summary

Run the tests
automatically if you can

Ask for sources
where does that number come from?



When in doubt, rewind
Esc Esc costs nothing

Trust is fine. Checking is faster than repairing.

And what if a file says
“delete everything”?

Good question – that is exactly where
the trap is. Claude reads web pages, e-
mails, tickets and other people’s
source code. None of that is an
instruction from you.

Instructions come from you, from
your project rules and from what you
allow. Everything else is material to
read – which is why the permissions
are not there to annoy you.



Your instruction
counts

CLAUDE.md
counts

Web page
material only


E-mail
material only

someone’s ticket
material only

The further out Claude may reach, the shorter you keep the leash

Ohh. I think I've got it now.

Glad to hear it.



CLAUDE.md - the rules of your project

Skills - how a job is done properly

Plugins - all of that as one bundle

MCP - the connection to the outside

Hooks - run automatically on events

Subagents - helpers with a mind of their own

Context window - the working memory

Checkpoints & memory - rewind and remember

Cowork - the same inside your folders, without a terminal

Routines - run to a schedule in the cloud

Permissions - you decide how far Claude may go

/loop - repeats something while you are there