


I get what `/loop` does now.

But Boris uses it to have Claude maintain his **whole** app – how does that work?



?

Auto-pacing

Crash fuzzer

Verification

Worktrees

Tuning

Dup unifier


Routine workshop

Ralph loop

Squad

Harness


That's called
Loop Engineering.
Come on, I'll
show you.



Hold on – I already know
`/loop`. So what is
Loop Engineering then?

`/loop` is just the building block: a
loop that repeats a task.

Loop Engineering is the craft around it
– making sure the loop can run
unattended for days without causing
harm.




`/loop`

the building block



Auto-pacing +
verification + tuning +
isolation



the craft around it



Give me an example of what that looks like in practice.

Boris Cherny runs several routines daily over his own apps, in a Slack channel.

A crash fuzzer taps around the app and fixes crashes, a dup unifier finds duplicate code, a dead-code remover cleans up – each routine with its own job.



Crash fuzzer

Dup unifier

Dead-code remover

Abstraction police





388 pull requests in a few weeks

180 of them merged after code review + human review

Do I have to tell the loop every time how often it should run?

Not necessarily. Without a time given, Claude picks the interval itself.

Short while things are actually happening – longer during quiet stretches. That's exactly what makes running for days practical: nobody has to keep adjusting an interval.



busy
short interval

quiet stretch
long interval

blocked
fallback wake-up





Claude decides for itself when the next run is worth it

This loop that just keeps planning, doing, checking – does that have a name too?

Yes – that's called a **Ralph-loop**: plan, execute, check, from the top.

Robust, but crude – without a brake it keeps burning tokens even after it's long since done. That's exactly what auto-pacing is for, and the verification that comes next.



↻ from the top, until done



Ralph loop

robust, but expensive in tokens without a brake

But if nobody's watching
- how am I supposed to
trust the loop?

That's exactly where Loop Engineering
stands or falls: the loop has to be able
to verify its own work **itself**,
end-to-end.

Run the tests, have a second model
review it - automated code review and
security review - and only then offer
the change as a pull request.






if a check fails, there is no proposal - no human has to look at it beforehand

And what if the loop makes a mess of it at first?

Then you adjust the routine – or just ask Claude to improve it itself.

Sometimes one day is enough, sometimes it takes several attempts before the routine sits reliably. That's expected, not a failure.





Claude gets feedback and tunes its own routine

Is that tuning from before systematic over time, or is it just guesswork?

The systematic version is called **hill climbing**: you measure quality with a fixed checklist – an eval – and adjust the routine on purpose.

No guesswork: every change gets tested against the eval, only what actually scores better stays in.





↗ step by step, uphill

keep only what the eval confirms – that's hill climbing

Don't several loops in the same folder get in each other's way?

No – each routine works in its own copy of the project, a **worktree**.

That way parallel runs don't interfere with each other, and a failed attempt doesn't drag the main branch down with it.



Main project



Worktree A

Worktree B



Worktree C

each loop on its own, nothing overwrites anything else

The worktrees from before – do they each run as their own, separate Claude?

Mostly yes: each routine its own Claude – many of them at once is called a **fleet**.


But sometimes several Claudes work **together** on a single task, with roles split up – that's a **squad**: one plans, one builds, one reviews.



Auto-pacing, verification, tuning, isolation – does this whole scaffolding have a name too?

Yes – the sum of tools, permissions, memory, and orchestration around the model is called a **harness**.

Loop Engineering is exactly that: building a harness Claude can work in, unattended, for days.





all of it together: the **harness**

So when is all this craft actually worth it – isn't /loop enough most of the time?


For a single afternoon, /loop really is completely enough.

Loop Engineering pays off once several such routines are meant to run alongside each other permanently – then individual loops turn into a small workshop that maintains itself.



Now Boris' Slack channel finally makes sense.

That's exactly the point.



/loop - the building block, repeats a task

Ralph loop - plan, execute, check, from the top

Auto-pacing - Claude picks the interval itself

Verification - tests, review, only then a proposal

Hill climbing - tuning with an eval instead of guesswork

Worktree isolation - each loop on its own

Squad / fleet - split roles, or many in parallel

Harness - tools, permissions, memory, orchestration

Loop Engineering - all of it as a workshop in permanent operation